

Machine Learning 10-601

Tom M. Mitchell
Machine Learning Department
Carnegie Mellon University

April 1, 2019

Today:

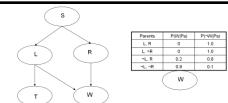
- Inference in graphical models
- Learning graphical models

Readings:

- Bishop chapter 8

Bayesian Networks Definition

A Bayes network represents the joint probability distribution over a collection of random variables



Parents	Prob(No)	Prob(Yes)
L	0.2	0.8
L+R	0.9	0.1
~L	0.2	0.8
~L+R	0.1	0.9
~L,R	0.3	0.7
~L,R,W	0.1	0.9

A Bayes network is a directed acyclic graph and a set of conditional probability distributions (CPD's)

- Each node denotes a random variable
- Edges denote dependencies
- For each node X_i its CPD defines $P(X_i | Pa(X_i))$
- The joint distribution over all variables is defined to be

$$P(X_1 \dots X_n) = \prod_i P(X_i | Pa(X_i))$$

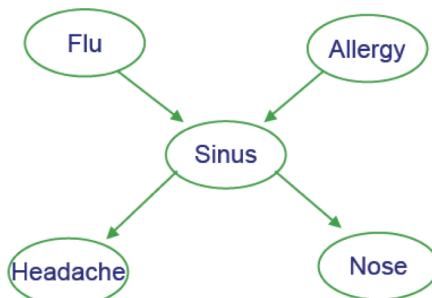
$Pa(X)$ = immediate parents of X in the graph

Inference in Bayes Nets

- In general, intractable (NP-complete)
- For certain cases, tractable

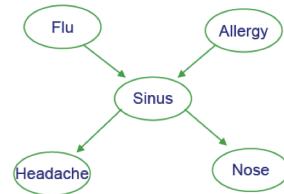
Example

- Flu and Allergies both cause Sinus problems
- Sinus problems cause Headaches and runny Nose



Prob. of joint assignment: easy

Suppose we are interested in joint assignment $\langle F=f, A=a, S=s, H=h, N=n \rangle$

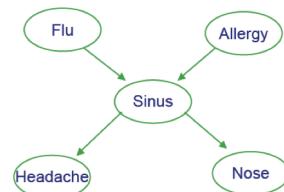


What is $P(f,a,s,h,n)$?

let's use $p(a,b)$ as shorthand for $p(A=a, B=b)$

Marginal probabilities $P(X_i)$: not so easy

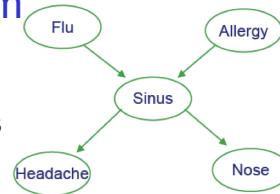
- How do we calculate $P(N=n)$?



let's use $p(a,b)$ as shorthand for $p(A=a, B=b)$

Generating a random sample from joint distribution: easy

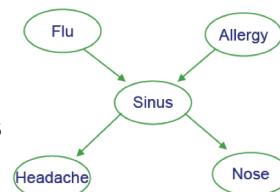
How can we generate random samples drawn according to $P(F,A,S,H,N)$?



let's use $p(a,b)$ as shorthand for $p(A=a, B=b)$

Generating a sample from joint distribution: easy

How can we generate random samples drawn according to $P(F,A,S,H,N)$?



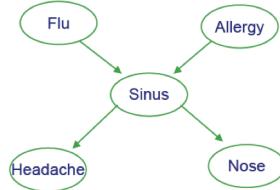
To generate a random sample for roots of network (F or A):

1. let $\theta = P(F=1)$ # look up from CPD
2. r = random number drawn uniformly between 0 and 1
3. if $r < \theta$ then output 1, else 0

let's use $p(a,b)$ as shorthand for $p(A=a, B=b)$

Generating a sample from joint distribution: easy

How can we generate random samples drawn according to $P(F,A,S,H,N)$?



To generate a random sample for roots of network (F or A):

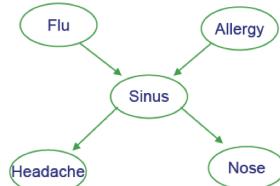
1. let $\theta = P(F=1)$ # look up from CPD
2. r = random number drawn uniformly between 0 and 1
3. if $r < \theta$ then output 1, else 0

To generate a random sample for S, given F,A:

1. let $\theta = P(S=1|F=f,A=a)$ # look up from CPD
2. r = random number drawn uniformly between 0 and 1
3. if $r < \theta$ then output 1, else 0

Generating a sample from joint distribution: easy

Note we can estimate marginals like $P(N=n)$ by generating many samples from joint distribution, then count the fraction of samples for which $N=n$



Similarly, for anything else we care about, calculate its maximum likelihood estimate from the sample
 $P(F=1|H=1, N=0)$

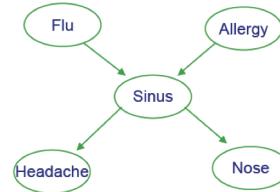
→ weak but general method for estimating any probability term...

Generating a sample from joint distribution: easy

We can easily sample $P(F,A,S,H,N)$

Can we use this to get $P(F,A,S,H | N)$?

Directly sample $P(F,A,S,H | N)$?



Gibbs Sampling:

Goal: Directly sample conditional distributions

$$P(X_1, \dots, X_n | X_{n+1}, \dots, X_m)$$

Approach:

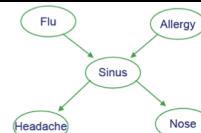
- start with the fixed observed X_{n+1}, \dots, X_m
plus arbitrary initial values for unobserved $X_1^{(0)}, \dots, X_n^{(0)}$
- iterate for $s=0$ to a big number:

$$X_1^{s+1} \sim P(X_1 | X_2^s, X_3^s, \dots, X_n^s, X_{n+1}, \dots, X_m)$$

$$X_2^{s+1} \sim P(X_2 | X_1^{s+1}, X_3^s, \dots, X_n^s, X_{n+1}, \dots, X_m)$$

...

$$X_n^{s+1} \sim P(X_n | X_1^{s+1}, X_2^{s+1}, \dots, X_{n-1}^{s+1}, X_{n+1}, \dots, X_m)$$



Eventually (after burn-in), the collection of samples will constitute a sample of the true $P(X_1, \dots, X_n | X_{n+1}, \dots, X_m)$

* but often use every 100th sample, since iters not independent

Gibbs Sampling:

Approach:

- start with arbitrary initial values for $X_1^{(0)}, \dots, X_n^{(0)}$
(and observed X_{n+1}, \dots, X_m)
- iterate for $s=0$ to a big number:

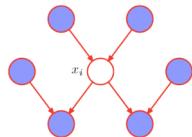
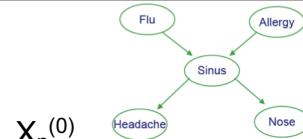
$$X_1^{s+1} \sim P(X_1 | X_2^s, X_3^s, \dots, X_n^s, X_{n+1}, \dots, X_m)$$

$$X_2^{s+1} \sim P(X_2 | X_1^{s+1}, X_3^s, \dots, X_n^s, X_{n+1}, \dots, X_m)$$

...

$$X_n^{s+1} \sim P(X_n | X_1^{s+1}, X_2^{s+1}, \dots, X_{n-1}^{s+1}, X_{n+1}, \dots, X_m)$$

Only need Markov Blanket at each step!

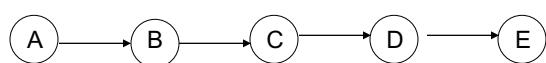


Gibbs is special case of Markov Chain Monte Carlo method

Prob. of marginals: not so easy

But sometimes the structure of the network allows us to be clever → avoid exponential work

e.g., chain



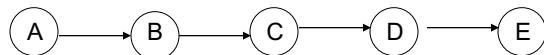
what is $P(C=1 | B=b, D=d)$?

what is $P(C=1)$?

Variable Elimination example

But sometimes the structure of the network allows us to be clever → avoid exponential work

eg., chain



what is $P(C=1)$?

Inference in Bayes Nets

- In general, intractable (NP-complete)
- For certain cases, tractable
 - Assigning probability to fully observed set of variables
 - Or if just one variable unobserved
 - Or for singly connected graphs (ie., no undirected loops)
 - Variable elimination
- Can often use Monte Carlo methods
 - Generate many samples, then count up the results
 - Gibbs sampling (example of Markov Chain Monte Carlo)
- Many other approaches
 - Variational methods for tractable approximate solutions
 - Junction tree, Belief propagation, ...

see Graphical Models course 10-708

Learning Bayes Nets from Data

Learning of Bayes Nets

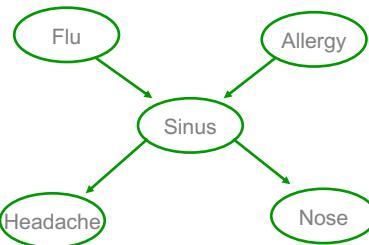
- Four categories of learning problems
 - Graph structure may be known/unknown
 - Variable values may be fully observed / partly unobserved
- Easy case: learn parameters when graph structure is *known*, and training data is *fully observed*
- Interesting case: graph *known*, data *partly observed*
- Gruesome case: graph structure *unknown*, data *partly unobserved*

Learning CPTs from Fully Observed Data

- Example: Consider learning the parameter

$$\theta_{s|ij} \equiv P(S = 1 | F = i, A = j)$$

- MLE (Max Likelihood Estimate) is



$$\theta_{s|ij} = \frac{\sum_{k=1}^K \delta(f_k = i, a_k = j, s_k = 1)}{\sum_{k=1}^K \delta(f_k = i, a_k = j)}$$

kth training example

$\delta(X) = 1$ if X is true
0 otherwise

- Remember why?

let's use a_k to represent value of A on the kth example

MLE estimate of $\theta_{s|ij}$ from fully observed data

- Maximum likelihood estimate

$$\theta \leftarrow \arg \max_{\theta} \log P(\text{data}|\theta)$$

- Our case:

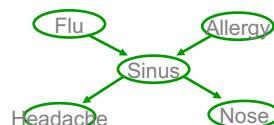
$$P(\text{data}|\theta) = \prod_{k=1}^K P(f_k, a_k, s_k, h_k, n_k)$$

$$P(\text{data}|\theta) = \prod_{k=1}^K P(f_k)P(a_k)P(s_k|f_k a_k)P(h_k|s_k)P(n_k|s_k)$$

$$\log P(\text{data}|\theta) = \sum_{k=1}^K \log P(f_k) + \log P(a_k) + \log P(s_k|f_k a_k) + \log P(h_k|s_k) + \log P(n_k|s_k)$$

$$\frac{\partial \log P(\text{data}|\theta)}{\partial \theta_{s|ij}} = \sum_{k=1}^K \frac{\partial \log P(s_k|f_k a_k)}{\partial \theta_{s|ij}}$$

$$\theta_{s|ij} = \frac{\sum_{k=1}^K \delta(f_k = i, a_k = j, s_k = 1)}{\sum_{k=1}^K \delta(f_k = i, a_k = j)}$$



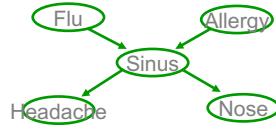
let's use a_k to represent value of A on the kth example

MLE for $\theta_{s|ij} = P(S = 1|F = i, A = j)$ from fully observed data

- Maximum likelihood estimate

$$\theta \leftarrow \arg \max_{\theta} \log P(\text{data}|\theta)$$

$$\theta_{s|ij} = \frac{\sum_{k=1}^K \delta(f_k = i, a_k = j, s_k = 1)}{\sum_{k=1}^K \delta(f_k = i, a_k = j)}$$



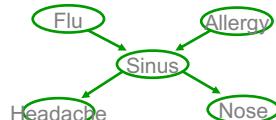
like flipping coin $\sum_{k=1}^K \delta(f_k = i, a_k = j)$ times to see
how often $s_k = 1$

MAP for $\theta_{s|ij} = P(S = 1|F = i, A = j)$ from fully observed data

- Maximum likelihood estimate

$$\theta \leftarrow \arg \max_{\theta} \log P(\text{data}|\theta)$$

$$\theta_{s|ij} = \frac{\sum_{k=1}^K \delta(f_k=i, a_k=j, s_k=1)}{\sum_{k=1}^K \delta(f_k=i, a_k=j)}$$



- MAP estimate

$$\theta \leftarrow \arg \max_{\theta} \log P(\theta|\text{data}) = \arg \max_{\theta} \log [P(\text{data}|\theta)P(\theta)]$$

If assume prior $P(\theta_{s|ij}) = \text{Beta}(\beta_1, \beta_0) = \frac{1}{B(\beta_1, \beta_0)} \theta_{s|ij}^{\beta_1-1} (1 - \theta_{s|ij})^{\beta_0-1}$

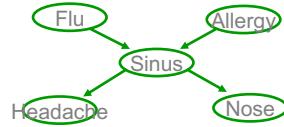
$$\theta_{s|ij} = \frac{(\beta_1-1)+\sum_{k=1}^K \delta(f_k=i, a_k=j, s_k=1)}{(\beta_1-1)+(\beta_0-1)+\sum_{k=1}^K \delta(f_k=i, a_k=j)}$$

like coin flipping, including hallucinated examples

Estimate θ from partly observed data

- What if FAHN observed, but not S?
- Can't calculate MLE

$$\theta \leftarrow \arg \max_{\theta} \log \prod_k P(f_k, a_k, s_k, h_k, n_k | \theta)$$



- Let X be all *observed* variable values (over all examples)
- Let Z be all *unobserved* variable values

- Can't calculate MLE:

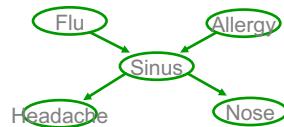
$$\theta \leftarrow \arg \max_{\theta} \log P(X, Z | \theta)$$

- WHAT TO DO?

Estimate θ from partly observed data

- What if FAHN observed, but not S?
- Can't calculate MLE

$$\theta \leftarrow \arg \max_{\theta} \log \prod_k P(f_k, a_k, s_k, h_k, n_k | \theta)$$



- Let X be all *observed* variable values (over all examples)
- Let Z be all *unobserved* variable values
- Can't calculate MLE:

$$\theta \leftarrow \arg \max_{\theta} \log P(X, Z | \theta)$$

- EM seeks* the estimate:

$$\theta \leftarrow \arg \max_{\theta} E_{Z|X,\theta} [\log P(X, Z | \theta)]$$

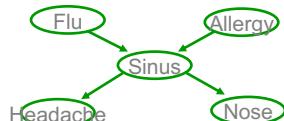
* EM guaranteed to find local maximum

Expected value

$$E_{P(X)}[f(X)] = \sum_x P(X = x)f(x)$$

- EM seeks estimate:

$$\theta \leftarrow \arg \max_{\theta} E_{Z|X,\theta}[\log P(X, Z|\theta)]$$



- here, observed $X=\{F,A,H,N\}$, unobserved $Z=\{S\}$

$$\log P(X, Z|\theta) = \sum_{k=1}^K \log P(f_k) + \log P(a_k) + \log P(s_k|f_k a_k) + \log P(h_k|s_k) + \log P(n_k|s_k)$$

$$E_{P(Z|X,\theta)} \log P(X, Z|\theta)$$

$$= \sum_{k=1}^K \sum_{i=0}^1 P(s_k = i | f_k, a_k, h_k, n_k) [\log P(f_k) + \log P(a_k) + \log P(s_k | f_k a_k) + \log P(h_k | s_k) + \log P(n_k | s_k)]$$

let's use a_k to represent value of A on the k th example

EM Algorithm - Informally



EM is a general procedure for learning from partly observed data

Given observed variables X, unobserved Z ($X=\{F,A,H,N\}$, $Z=\{S\}$)

Begin with arbitrary choice for parameters θ

Iterate until convergence:

- E Step: use X, θ to estimate the unobserved Z values
- M Step: use X values and estimated Z values to derive a better θ

Guaranteed to find local maximum.

Each iteration increases $E_{P(Z|X,\theta)}[\log P(X, Z|\theta')]$

EM Algorithm - Precisely

EM is a general procedure for learning from partly observed data

Given observed variables X, unobserved Z ($X=\{F,A,H,N\}$, $Z=\{S\}$)

Define $Q(\theta'|\theta) = E_{P(Z|X,\theta)}[\log P(X, Z|\theta')]$

\uparrow current \nwarrow M step new

Iterate until convergence:

- E Step: Use X and current θ to calculate $P(Z|X,\theta)$
- M Step: Replace current θ by

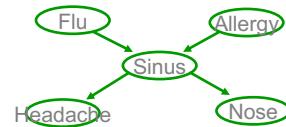
$$\theta \leftarrow \arg \max_{\theta'} Q(\theta'|\theta)$$

Guaranteed to find local maximum.

Each iteration increases $E_{P(Z|X,\theta)}[\log P(X, Z|\theta')]$

E Step: Use X, θ , to Calculate $P(Z|X, \theta)$

observed $X = \{F, A, H, N\}$,
unobserved $Z = \{S\}$



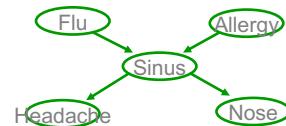
How? Bayes net inference problem.

$$P(S_k = 1 | f_k a_k h_k n_k, \theta) =$$

let's use a_k to represent value of A on the k th example

E Step: Use X, θ , to Calculate $P(Z|X, \theta)$

observed $X = \{F, A, H, N\}$,
unobserved $Z = \{S\}$



How? Bayes net inference problem.

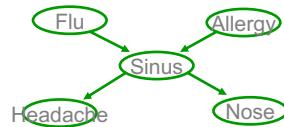
$$P(S_k = 1 | f_k a_k h_k n_k, \theta) =$$

$$P(S_k = 1 | f_k a_k h_k n_k, \theta) = \frac{P(S_k = 1, f_k a_k h_k n_k | \theta)}{P(S_k = 1, f_k a_k h_k n_k | \theta) + P(S_k = 0, f_k a_k h_k n_k | \theta)}$$

let's use a_k to represent value of A on the k th example

EM and estimating $\theta_{s|ij}$

observed X = {F,A,H,N}, unobserved Z={S}



E step: Calculate $P(Z_k|X_k; \theta)$ for each training example, k

$$P(S_k = 1|f_k a_k h_k n_k, \theta) = E[s_k] = \frac{P(S_k = 1, f_k a_k h_k n_k | \theta)}{P(S_k = 1, f_k a_k h_k n_k | \theta) + P(S_k = 0, f_k a_k h_k n_k | \theta)}$$

M step: update all relevant parameters. For example:

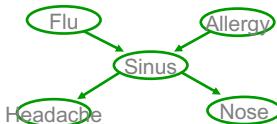
$$\theta_{s|ij} \leftarrow \frac{\sum_{k=1}^K \delta(f_k = i, a_k = j) E[s_k]}{\sum_{k=1}^K \delta(f_k = i, a_k = j)}$$

Recall MLE was: $\theta_{s|ij} = \frac{\sum_{k=1}^K \delta(f_k = i, a_k = j, s_k = 1)}{\sum_{k=1}^K \delta(f_k = i, a_k = j)}$

EM and estimating θ

More generally,

Given observed set X, unobserved set Z of boolean values



E step: Calculate for each training example, k

the expected value of each unobserved variable in each training example

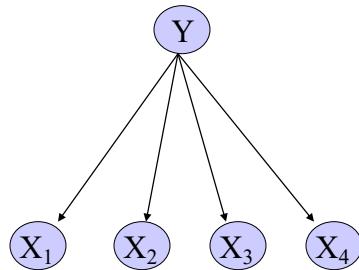
M step:

Calculate θ similar to MLE estimates, but replacing each count by its expected count

$$\delta(Z = 1) \rightarrow E_{Z|X,\theta}[Z] \quad \delta(Z = 0) \rightarrow (1 - E_{Z|X,\theta}[Z])$$

Using Unlabeled Data to Help Train Naïve Bayes Classifier

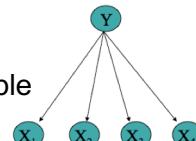
Learn $P(Y|X)$



Y	X1	X2	X3	X4
1	0	0	1	1
0	0	1	0	0
0	0	0	1	0
?	0	1	1	0
?	0	1	0	1

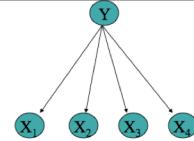
E step: Calculate for each training example, k

the expected value of each unobserved variable



Y	X1	X2	X3	X4
1	0	0	1	1
0	0	1	0	0
0	0	0	1	0
?	0	1	1	0
?	0	1	0	1

EM and estimating θ



Given observed set X, unobserved set Y of boolean values

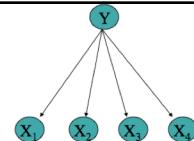
E step: Calculate for each training example, k
the expected value of each unobserved variable Y

$$E_{P(Y|X_1 \dots X_N)}[y(k)] = P(y(k) = 1|x_1(k), \dots x_N(k); \theta) = \frac{P(y(k) = 1) \prod_i P(x_i(k)|y(k) = 1)}{\sum_{j=0}^1 P(y(k) = j) \prod_i P(x_i(k)|y(k) = j)}$$

M step: Calculate estimates similar to MLE, but
replacing each count by its expected count

let's use $y(k)$ to indicate value of Y on kth example

EM and estimating θ



Given observed set X, unobserved set Y of boolean values

E step: Calculate for each training example, k
the expected value of each unobserved variable Y

$$E_{P(Y|X_1 \dots X_N)}[y(k)] = P(y(k) = 1|x_1(k), \dots x_N(k); \theta) = \frac{P(y(k) = 1) \prod_i P(x_i(k)|y(k) = 1)}{\sum_{j=0}^1 P(y(k) = j) \prod_i P(x_i(k)|y(k) = j)}$$

M step: Calculate estimates similar to MLE, but
replacing each count by its expected count

$$\theta_{ij|m} = \hat{P}(X_i = j|Y = m) = \frac{\sum_k P(y(k) = m|x_1(k) \dots x_N(k)) \delta(x_i(k) = j)}{\sum_k P(y(k) = m|x_1(k) \dots x_N(k))}$$

$$\text{MLE would be: } \hat{P}(X_i = j|Y = m) = \frac{\sum_k \delta((y(k) = m) \wedge (x_i(k) = j))}{\sum_k \delta(y(k) = m)}$$

-
- **Inputs:** Collections \mathcal{D}^l of labeled documents and \mathcal{D}^u of unlabeled documents.
 - Build an initial naive Bayes classifier, $\hat{\theta}$, from the labeled documents, \mathcal{D}^l , only. Use maximum a posteriori parameter estimation to find $\hat{\theta} = \arg \max_{\theta} P(\mathcal{D}|\theta)P(\theta)$ (see Equations 5 and 6).
 - Loop while classifier parameters improve, as measured by the change in $l_c(\theta|\mathcal{D}; \mathbf{z})$ (the complete log probability of the labeled and unlabeled data)
 - **(E-step)** Use the current classifier, $\hat{\theta}$, to estimate component membership of each unlabeled document, *i.e.*, the probability that each mixture component (and class) generated each document, $P(c_j|d_i; \hat{\theta})$ (see Equation 7).
 - **(M-step)** Re-estimate the classifier, $\hat{\theta}$, given the estimated component membership of each document. Use maximum a posteriori parameter estimation to find $\hat{\theta} = \arg \max_{\theta} P(\mathcal{D}|\theta)P(\theta)$ (see Equations 5 and 6).
 - **Output:** A classifier, $\hat{\theta}$, that takes an unlabeled document and predicts a class label.

From [Nigam et al., 2000]



Experimental Evaluation

- **Newsgroup postings**
 - 20 newsgroups, 1000/group
- **Web page classification**
 - student, faculty, course, project
 - 4199 web pages
- **Reuters newswire articles**
 - 12,902 articles
 - 90 topics categories

20 Newsgroups

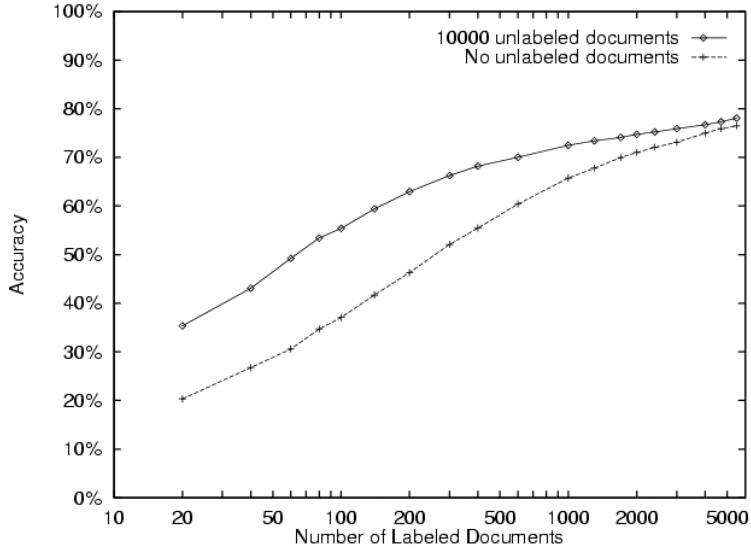


Table 3. Lists of the words most predictive of the `course` class in the WebKB data set, as they change over iterations of EM for a specific trial. By the second iteration of EM, many common `course`-related words appear. The symbol D indicates an arbitrary digit.

Iteration 0		Iteration 1	Iteration 2
intelligence	word w ranked by $P(w Y=\text{course})$ $/P(w Y \neq \text{course})$	DD	D
DD		D	DD
artificial		lecture	lecture
understanding		cc	cc
DDw		D^*	$DD:DD$
dist		$DD:DD$	due
identical		handout	D^*
rus		due	homework
arrange		problem	assignment
games		set	handout
dartmouth		tay	set
natural		$DDam$	hw
cognitive		yurttas	exam
logic		homework	problem
proving	Using one labeled example per class	kfoury	$DDam$
prolog		sec	postscript
knowledge		postscript	solution
human		exam	quiz
representation		solution	chapter
field		assaf	ascii

What you should know about EM

- For learning from partly unobserved data
- MLE of $\theta = \arg \max_{\theta} \log P(\text{data}|\theta)$
- EM estimate: $\theta = \arg \max_{\theta} E_{Z|X,\theta}[\log P(X, Z|\theta)]$
Where X is observed part of data, Z is unobserved
- EM for training Bayes networks
- Recall EM for Gaussian Mixture Models
- Can also derive your own EM algorithm for your own problem
 - write out expression for $E_{Z|X,\theta}[\log P(X, Z|\theta)]$
 - E step: for each training example X^k , calculate $P(Z^k | X^k, \theta)$
 - M step: chose new θ to maximize $E_{Z|X,\theta}[\log P(X, Z|\theta)]$

Learning Bayes Net Structure

How can we learn Bayes Net graph structure?

In general case, open problem

- can require lots of data (else high risk of overfitting)
- can use Bayesian priors, or other kinds of prior assumptions about graph structure to constrain search

One key result:

- Chow-Liu algorithm: finds “best” tree-structured network
- What’s best?
 - suppose $P(\mathbf{X})$ is true distribution, $T(\mathbf{X})$ is our tree-structured network, where $\mathbf{X} = \langle X_1, \dots, X_n \rangle$
 - Chow-Liu minimizes Kullback-Leibler divergence:

$$KL(P(\mathbf{X}) \parallel T(\mathbf{X})) \equiv \sum_k P(\mathbf{X} = k) \log \frac{P(\mathbf{X} = k)}{T(\mathbf{X} = k)}$$

Kullback-Leibler Divergence

- $KL(P(\mathbf{X}) \parallel T(\mathbf{X}))$ is a measure of the difference between distribution $P(\mathbf{X})$ and $T(\mathbf{X})$

$$KL(P(\mathbf{X}) \parallel T(\mathbf{X})) \equiv \sum_k P(\mathbf{X} = k) \log \frac{P(\mathbf{X} = k)}{T(\mathbf{X} = k)}$$

- It is assymetric, always greater or equal to 0
- It is 0 iff $P(\mathbf{X})=T(\mathbf{X})$

$$\begin{aligned} KL(P(X) \parallel T(X)) &= \sum_k P(X = k) \log P(X = k) - \sum_k P(X = k) \log T(X = k) \\ &= -H(P) + H(P, T) \end{aligned}$$

where cross entropy $H(P, T) = \sum_k -P(X = k) \log T(X = k)$

Chow-Liu Algorithm

Key result: To minimize $KL(P \parallel T)$ over possible tree networks T representing true P , it suffices to find the tree network T that maximizes the sum of mutual informations over its edges

Mutual information for an edge between variable A and B:

$$I(A, B) = \sum_a \sum_b P(a, b) \log \frac{P(a, b)}{P(a)P(b)}$$

This works because for tree networks with nodes $\mathbf{X} \equiv \langle X_1 \dots X_n \rangle$

$$\begin{aligned} KL(P(\mathbf{X}) \parallel T(\mathbf{X})) &\equiv \sum_k P(\mathbf{X} = k) \log \frac{P(\mathbf{X} = k)}{T(\mathbf{X} = k)} \\ &= -\sum_i I(X_i, Pa(X_i)) + \sum_i H(X_i) - H(X_1 \dots X_n) \end{aligned}$$

Chow-Liu Algorithm

1. for each pair of variables A,B, use data to estimate $P(A,B)$, $P(A)$, and $P(B)$

2. for each pair A, B calculate mutual information

$$I(A, B) = \sum_a \sum_b P(a, b) \log \frac{P(a, b)}{P(a)P(b)}$$

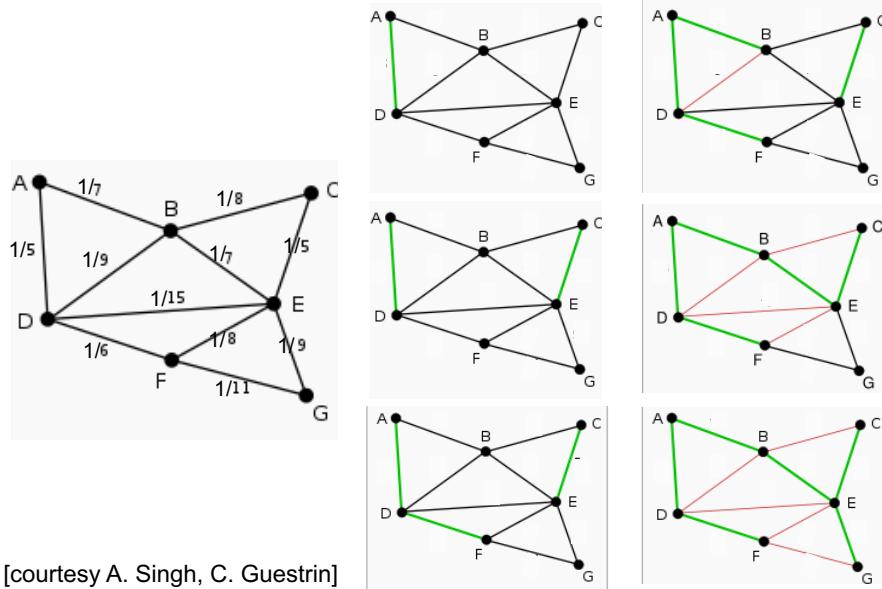
3. calculate the maximum spanning tree over the set of variables, using edge weights $I(A, B)$
(given N vars, this costs only $O(N^2)$ time)

4. add arrows to edges to form a directed-acyclic graph

5. learn the CPD's for this graph

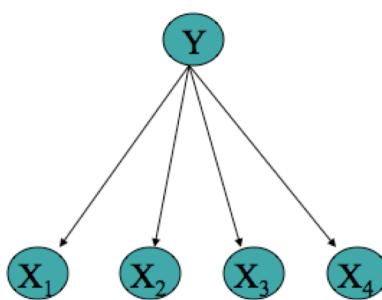
Chow-Liu algorithm example

Greedy Algorithm to find Max-Spanning Tree



Tree Augmented Naïve Bayes

[Nir Friedman et al., 1997]



Bayes Nets – What You Should Know

- Representation
 - Bayes nets represent joint distribution as a DAG + Conditional Distributions
- Inference
 - NP-hard in general
 - For some graphs, closed form inference is feasible
 - Approximate methods too, e.g., Monte Carlo methods, ...
- Learning
 - Easy for known graph, fully observed data (MLE's, MAP est.)
 - EM for partly observed data, known graph
 - Learning graph structure: Chow-Liu for tree-structured networks
 - Hardest when graph unknown, data incompletely observed