

## 10-701 Introduction to Machine Learning (PhD) Lecture 6: Logistic Regression

Leila Wehbe  
Carnegie Mellon University  
Machine Learning Department

Slides based on Tom Mitchell's  
10-701 Spring 2016 material

### Naïve Bayes Algorithm – discrete $X_i$

- Train Naïve Bayes (examples)
  - for each\* value  $y_k$   
estimate  $\pi_k \equiv P(Y = y_k)$
  - for each\* value  $x_{ij}$  of each attribute  $X_i$   
estimate  $\theta_{ijk} \equiv P(X_i = x_{ij}|Y = y_k)$
- Classify ( $X^{new}$ )

$$Y^{new} \leftarrow \arg \max_{y_k} P(Y = y_k) \prod_i P(X_i^{new}|Y = y_k)$$

$$Y^{new} \leftarrow \arg \max_{y_k} \pi_k \prod_i \theta_{ijk}$$

\* probabilities must sum to 1, so need estimate only v-1 of these, where v is the number of values, which is 2 in the binary case

### Last time: Naïve Bayes in a Nutshell

Bayes rule:

$$P(Y = y_k|X_1 \dots X_n) = \frac{P(Y = y_k)P(X_1 \dots X_n|Y = y_k)}{\sum_j P(Y = y_j)P(X_1 \dots X_n|Y = y_j)}$$

Assuming conditional independence among  $X_i$ 's:

$$P(Y = y_k|X_1 \dots X_n) = \frac{P(Y = y_k) \prod_i P(X_i|Y = y_k)}{\sum_j P(Y = y_j) \prod_i P(X_i|Y = y_j)} \quad \begin{matrix} \text{(estimate} \\ \text{in} \\ \text{training)} \end{matrix}$$

So, to pick most probable Y for  $X^{new} = (X_1, \dots, X_n)$

$$Y^{new} \leftarrow \arg \max_{y_k} P(Y = y_k) \prod_i P(X_i^{new}|Y = y_k) \quad \begin{matrix} \text{(testing)} \end{matrix}$$

### MAP estimates for bag of words

Map estimate for multinomial

$$\theta_{jk} = \frac{\alpha_{jk} + \beta_{jk} - 1}{\sum_m \alpha_{mk} + \sum_m \beta_{mk} - 1}$$

# seen "aardvark"      # hallucinated "aardvark"  
# seen words      # hallucinated words

What  $\beta$ 's should we choose?

Missing data?

## What if we have continuous $X_i$ ?

Gaussian Naïve Bayes (GNB): assume

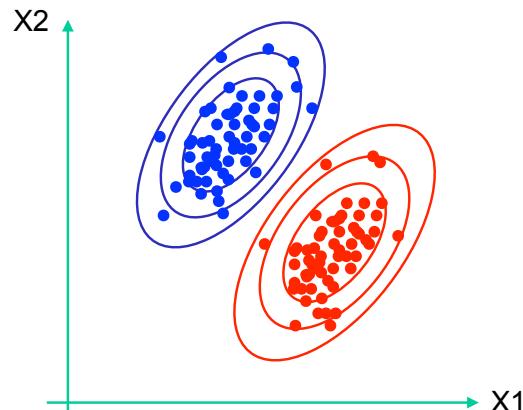
$$p(X_i = x|Y = y_k) = \frac{1}{\sqrt{2\pi\sigma_{ik}^2}} e^{-\frac{1}{2}\left(\frac{x-\mu_{ik}}{\sigma_{ik}}\right)^2}$$

Sometimes assume variance

- is independent of  $Y$  (i.e.,  $\sigma_i$ ),
- or independent of  $X_i$  (i.e.,  $\sigma_k$ )
- or both (i.e.,  $\sigma$ )

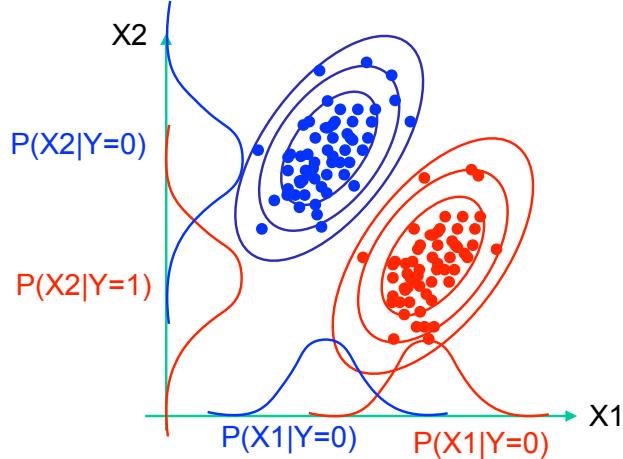
## Gaussian Naïve Bayes – Big Picture

$$Y^{new} \leftarrow \arg \max_{y \in \{0,1\}} P(Y = y) \prod_i P(X_i^{new}|Y = y) \quad \text{assume } P(Y=1) = 0.5$$



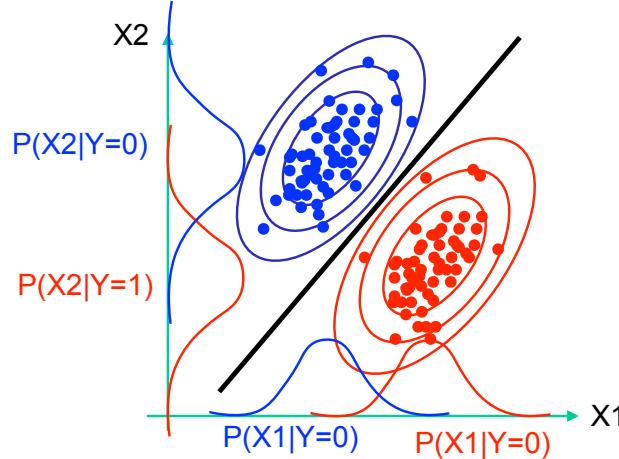
## Gaussian Naïve Bayes – Big Picture

$$Y^{new} \leftarrow \arg \max_{y \in \{0,1\}} P(Y = y) \prod_i P(X_i^{new}|Y = y) \quad \text{assume } P(Y=1) = 0.5$$



## Gaussian Naïve Bayes – Big Picture

$$Y^{new} \leftarrow \arg \max_{y \in \{0,1\}} P(Y = y) \prod_i P(X_i^{new}|Y = y) \quad \text{assume } P(Y=1) = 0.5$$



## Logistic Regression

Idea:

- Naïve Bayes allows computing  $P(Y|X)$  by learning  $P(Y)$  and  $P(X|Y)$
- Why not learn  $P(Y|X)$  directly?

- Consider learning  $f: X \rightarrow Y$ , where
  - $X$  is a vector of real-valued features, ( $X_1 \dots X_n$ )
  - $Y$  is boolean
  - assume all  $X_i$  are conditionally independent given  $Y$
  - model  $P(X_i | Y = y_k)$  as Gaussian  $N(\mu_{ik}, \sigma_i)$
  - model  $P(Y)$  as Bernoulli ( $\pi$ )

- What does that imply about the form of  $P(Y|X)$ ?

$$P(Y = 1 | X = (X_1, \dots, X_n)) = \frac{1}{1 + \exp(w_0 + \sum_i w_i X_i)}$$

Derive form for  $P(Y|X)$  for Gaussian  $P(X_i | Y=y_k)$  assuming  $\sigma_{ik} = \sigma_i$

$$P(Y = 1 | X) = \frac{P(Y = 1)P(X|Y = 1)}{P(Y = 1)P(X|Y = 1) + P(Y = 0)P(X|Y = 0)}$$

Derive form for  $P(Y|X)$  for Gaussian  $P(X_i | Y=y_k)$  assuming  $\sigma_{ik} = \sigma_i$

$$\begin{aligned} P(Y = 1 | X) &= \frac{P(Y = 1)P(X|Y = 1)}{P(Y = 1)P(X|Y = 1) + P(Y = 0)P(X|Y = 0)} \\ &= \frac{1}{1 + \frac{P(Y=0)P(X|Y=0)}{P(Y=1)P(X|Y=1)}} \end{aligned}$$

Derive form for  $P(Y|X)$  for Gaussian  $P(X_i|Y=y_k)$  assuming  $\sigma_{ik} = \sigma_i$

$$\begin{aligned} P(Y=1|X) &= \frac{P(Y=1)P(X|Y=1)}{P(Y=1)P(X|Y=1) + P(Y=0)P(X|Y=0)} \\ &= \frac{1}{1 + \frac{P(Y=0)P(X|Y=0)}{P(Y=1)P(X|Y=1)}} \\ &= \frac{1}{1 + \exp(\ln \frac{P(Y=0)P(X|Y=0)}{P(Y=1)P(X|Y=1)})} \end{aligned}$$

Derive form for  $P(Y|X)$  for Gaussian  $P(X_i|Y=y_k)$  assuming  $\sigma_{ik} = \sigma_i$

$$\begin{aligned} P(Y=1|X) &= \frac{P(Y=1)P(X|Y=1)}{P(Y=1)P(X|Y=1) + P(Y=0)P(X|Y=0)} \\ &= \frac{1}{1 + \frac{P(Y=0)P(X|Y=0)}{P(Y=1)P(X|Y=1)}} \\ &= \frac{1}{1 + \exp(\ln \frac{P(Y=0)P(X|Y=0)}{P(Y=1)P(X|Y=1)})} \\ &= \frac{1}{1 + \exp((\ln \frac{1-\pi}{\pi}) + \boxed{\sum_i \ln \frac{P(X_i|Y=0)}{P(X_i|Y=1)}})} \end{aligned}$$

Derive form for  $P(Y|X)$  for Gaussian  $P(X_i|Y=y_k)$  assuming  $\sigma_{ik} = \sigma_i$

$$\begin{aligned} P(Y=1|X) &= \frac{P(Y=1)P(X|Y=1)}{P(Y=1)P(X|Y=1) + P(Y=0)P(X|Y=0)} \\ &= \frac{1}{1 + \frac{P(Y=0)P(X|Y=0)}{P(Y=1)P(X|Y=1)}} \\ &= \frac{1}{1 + \exp(\ln \frac{P(Y=0)P(X|Y=0)}{P(Y=1)P(X|Y=1)})} \\ &= \frac{1}{1 + \exp((\ln \frac{1-\pi}{\pi}) + \boxed{\sum_i \ln \frac{P(X_i|Y=0)}{P(X_i|Y=1)}})} \end{aligned}$$

$$\begin{aligned} P(X_i = x_i|Y = y_k) &= \frac{1}{\sigma_{ik}\sqrt{2\pi}} e^{\frac{-(x_i - \mu_{ik})^2}{2\sigma_{ik}^2}} \\ \ln P(X_i = x_i|Y = y_k) &= \frac{1}{\sigma_{ik}\sqrt{2\pi}} + \frac{-(x_i - \mu_{ik})^2}{2\sigma_{ik}^2} \\ \ln \frac{P(X_i = x_i|Y = 0)}{P(X_i = x_i|Y = 1)} &= \frac{1}{\sigma_{i0}\sqrt{2\pi}} + \frac{-(x_i - \mu_{i0})^2}{2\sigma_{i0}^2} - \frac{1}{\sigma_{i1}\sqrt{2\pi}} + \frac{-(x_i - \mu_{i1})^2}{2\sigma_{i1}^2} \\ &= \frac{1}{\sigma_{i0}\sqrt{2\pi}} - \frac{1}{\sigma_{i1}\sqrt{2\pi}} - \frac{x_i^2 - 2x_i\mu_{i0} + \mu_{i0}^2}{2\sigma_{i0}^2} + \frac{x_i^2 - 2x_i\mu_{i1} + \mu_{i1}^2}{2\sigma_{i1}^2} \end{aligned}$$

Now assume  
 $\sigma_{ik} = \sigma_i$

Derive form for  $P(Y|X)$  for Gaussian  $P(X_i|Y=y_k)$  assuming  $\sigma_{ik} = \sigma_i$

$$\begin{aligned} P(Y=1|X) &= \frac{P(Y=1)P(X|Y=1)}{P(Y=1)P(X|Y=1) + P(Y=0)P(X|Y=0)} \\ &= \frac{1}{1 + \frac{P(Y=0)P(X|Y=0)}{P(Y=1)P(X|Y=1)}} \\ &= \frac{1}{1 + \exp(\ln \frac{P(Y=0)P(X|Y=0)}{P(Y=1)P(X|Y=1)})} \\ &= \frac{1}{1 + \exp((\ln \frac{1-\pi}{\pi}) + \boxed{\sum_i \ln \frac{P(X_i|Y=0)}{P(X_i|Y=1)}})} \end{aligned}$$

Linear function!

$$P(Y=1|X) = \frac{1}{1 + \exp(w_0 + \sum_{i=1}^n w_i X_i)}$$

$$\sum_i \left( \frac{\mu_{i0} - \mu_{i1}}{\sigma_i^2} X_i + \frac{\mu_{i1}^2 - \mu_{i0}^2}{2\sigma_i^2} \right)$$

## Very convenient!

$$P(Y = 1|X = (X_1, \dots, X_n)) = \frac{1}{1 + \exp(w_0 + \sum_i w_i X_i)}$$

$$P(Y = 0|X = (X_1, \dots, X_n)) =$$

implies

$$\frac{P(Y = 0|X)}{P(Y = 1|X)} =$$

implies

$$\ln \frac{P(Y = 0|X)}{P(Y = 1|X)} =$$

## Very convenient!

$$P(Y = 1|X = (X_1, \dots, X_n)) = \frac{1}{1 + \exp(w_0 + \sum_i w_i X_i)}$$

$$P(Y = 0|X = (X_1, \dots, X_n)) = \frac{\exp(w_0 + \sum_i w_i X_i)}{1 + \exp(w_0 + \sum_i w_i X_i)}$$

implies

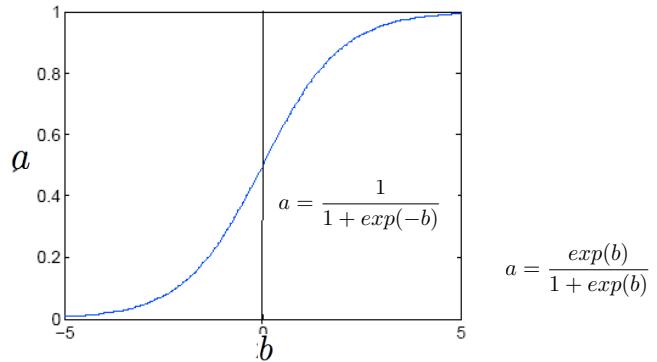
$$\frac{P(Y = 0|X)}{P(Y = 1|X)} = \exp(w_0 + \sum_i w_i X_i)$$

linear classification rule!

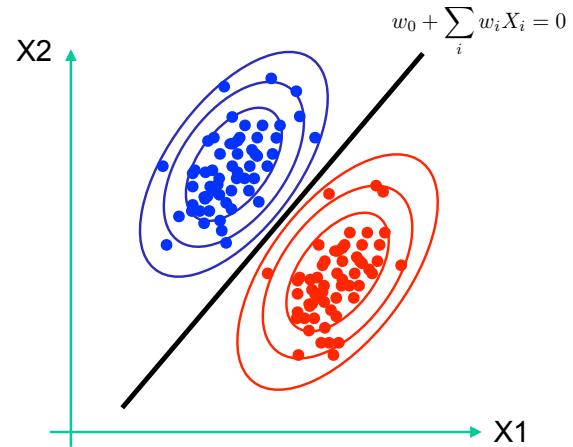
implies

$$\ln \frac{P(Y = 0|X)}{P(Y = 1|X)} = w_0 + \sum_i w_i X_i < 0$$

## Logistic function



$$P(Y = 1|X) = \frac{1}{1 + \exp(w_0 + \sum_{i=1}^n w_i X_i)}$$



## Logistic regression more generally

- Logistic regression when Y not boolean (but still discrete-valued).
- Now  $y \in \{y_1 \dots y_R\}$  : learn  $R-1$  sets of weights

$$\text{for } k < R \quad P(Y = y_k | X) = \frac{\exp(w_{k0} + \sum_{i=1}^n w_{ki}X_i)}{1 + \sum_{j=1}^{R-1} \exp(w_{j0} + \sum_{i=1}^n w_{ji}X_i)}$$

$$\text{for } k = R \quad P(Y = y_R | X) = \frac{1}{1 + \sum_{j=1}^{R-1} \exp(w_{j0} + \sum_{i=1}^n w_{ji}X_i)}$$

## Training Logistic Regression: MLE, MCLE

- we have L training examples:  $\{(X^1, Y^1), \dots, (X^L, Y^L)\}$
- maximum likelihood estimate for parameters W

$$\begin{aligned} W_{MLE} &= \arg \max_W P((X^1, Y^1), \dots, (X^L, Y^L) | W) \\ &= \arg \max_W \prod_l P(X^l, Y^l | W) \end{aligned}$$

- maximum conditional likelihood estimate

$$W_{M(C)LE} = \arg \max_W \prod_l P(Y^l | X^l, W)$$

## Training Logistic Regression: MCLE

- Choose parameters  $W = \langle w_0, \dots, w_n \rangle$  to maximize conditional likelihood of training data

where  $P(Y = 0 | X, W) = \frac{1}{1 + \exp(w_0 + \sum_i w_i X_i)}$

$$P(Y = 1 | X, W) = \frac{\exp(w_0 + \sum_i w_i X_i)}{1 + \exp(w_0 + \sum_i w_i X_i)}$$

- Training data D =  $\{(X^1, Y^1), \dots, (X^L, Y^L)\}$
- Data likelihood =  $\prod_l P(X^l, Y^l | W)$
- Data conditional likelihood =  $\prod_l P(Y^l | X^l, W)$

$$W_{MCLE} = \arg \max_W \prod_l P(Y^l | W, X^l)$$

## Expressing Conditional Log Likelihood

$$l(W) \equiv \ln \prod_l P(Y^l | X^l, W) = \sum_l \ln P(Y^l | X^l, W)$$

$$P(Y = 0 | X, W) = \frac{1}{1 + \exp(w_0 + \sum_i w_i X_i)}$$

$$P(Y = 1 | X, W) = \frac{\exp(w_0 + \sum_i w_i X_i)}{1 + \exp(w_0 + \sum_i w_i X_i)}$$

$$l(W) = \sum_l Y^l \ln P(Y^l = 1 | X^l, W) + (1 - Y^l) \ln P(Y^l = 0 | X^l, W)$$

## Expressing Conditional Log Likelihood

$$l(W) \equiv \ln \prod_l P(Y^l|X^l, W) = \sum_l \ln P(Y^l|X^l, W)$$

$$P(Y = 0|X, W) = \frac{1}{1 + \exp(w_0 + \sum_i w_i X_i)}$$

$$P(Y = 1|X, W) = \frac{\exp(w_0 + \sum_i w_i X_i)}{1 + \exp(w_0 + \sum_i w_i X_i)}$$

$$\begin{aligned} l(W) &= \sum_l Y^l \ln P(Y^l = 1|X^l, W) + (1 - Y^l) \ln P(Y^l = 0|X^l, W) \\ &= \sum_l Y^l \ln \frac{P(Y^l = 1|X^l, W)}{P(Y^l = 0|X^l, W)} + \ln P(Y^l = 0|X^l, W) \end{aligned}$$

## Expressing Conditional Log Likelihood

$$l(W) \equiv \ln \prod_l P(Y^l|X^l, W) = \sum_l \ln P(Y^l|X^l, W)$$

$$P(Y = 0|X, W) = \frac{1}{1 + \exp(w_0 + \sum_i w_i X_i)}$$

$$P(Y = 1|X, W) = \frac{\exp(w_0 + \sum_i w_i X_i)}{1 + \exp(w_0 + \sum_i w_i X_i)}$$

$$\begin{aligned} l(W) &= \sum_l Y^l \ln P(Y^l = 1|X^l, W) + (1 - Y^l) \ln P(Y^l = 0|X^l, W) \\ &= \sum_l Y^l \ln \frac{P(Y^l = 1|X^l, W)}{P(Y^l = 0|X^l, W)} + \ln P(Y^l = 0|X^l, W) \\ &= \sum_l Y^l (w_0 + \sum_i^n w_i X_i^l) - \ln(1 + \exp(w_0 + \sum_i^n w_i X_i^l)) \end{aligned}$$

## Maximizing Conditional Log Likelihood

$$P(Y = 0|X, W) = \frac{1}{1 + \exp(w_0 + \sum_i w_i X_i)}$$

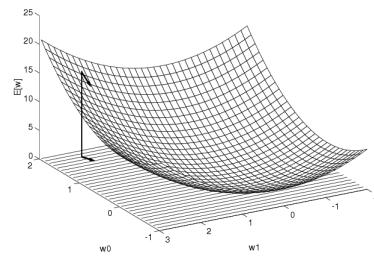
$$P(Y = 1|X, W) = \frac{\exp(w_0 + \sum_i w_i X_i)}{1 + \exp(w_0 + \sum_i w_i X_i)}$$

$$\begin{aligned} l(W) &\equiv \ln \prod_l P(Y^l|X^l, W) = \sum_l \ln P(Y^l|X^l, W) \\ &= \sum_l Y^l (w_0 + \sum_i^n w_i X_i^l) - \ln(1 + \exp(w_0 + \sum_i^n w_i X_i^l)) \end{aligned}$$

Good news:  $l(W)$  is concave function of  $W$

Bad news: no closed-form solution to maximize  $l(W)$

## Gradient Descent



Gradient

$$\nabla E[\vec{w}] \equiv \left[ \frac{\partial E}{\partial w_0}, \frac{\partial E}{\partial w_1}, \dots, \frac{\partial E}{\partial w_n} \right]$$

Training rule:

$$\Delta \vec{w} = -\eta \nabla E[\vec{w}]$$

i.e.,

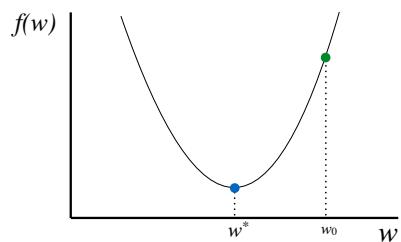
$$\Delta w_i = -\eta \frac{\partial E}{\partial w_i}$$

Update the vector of parameters

$$\mathbf{w} \leftarrow \mathbf{w} - \eta \nabla E_d(\mathbf{w})$$

## Gradient Descent

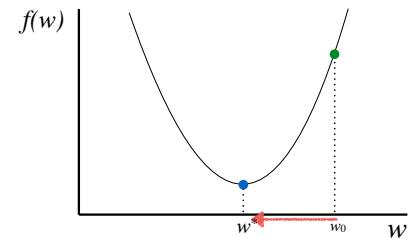
Start at a random point



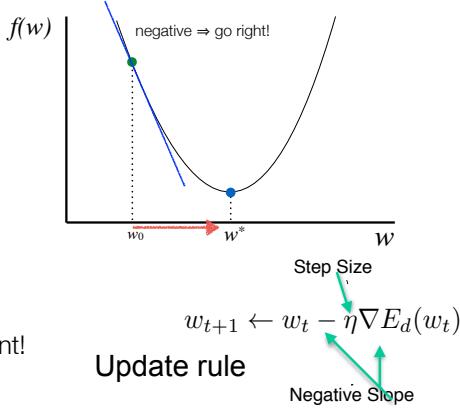
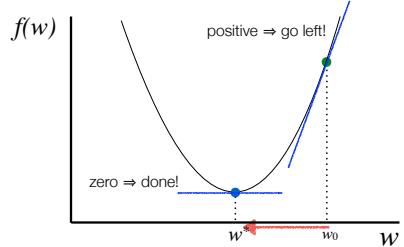
## Gradient Descent

Start at a random point

Determine a descent direction



## Choosing Descent Direction (1D)

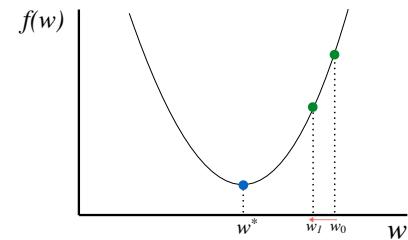


We can only move in two directions  
Negative slope is direction of descent!

## Gradient Descent

Start at a random point

Determine a descent direction  
Choose a step size  
Update



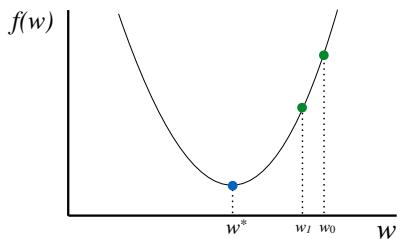
## Gradient Descent

Start at a random point

**Repeat**

- Determine a descent direction
- Choose a step size
- Update

**Until** stopping criterion is satisfied



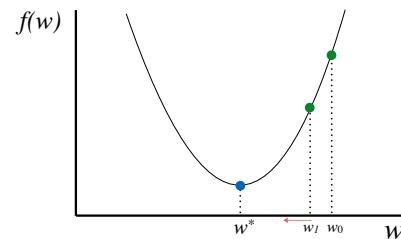
## Gradient Descent

Start at a random point

**Repeat**

- Determine a descent direction
- Choose a step size
- Update

**Until** stopping criterion is satisfied



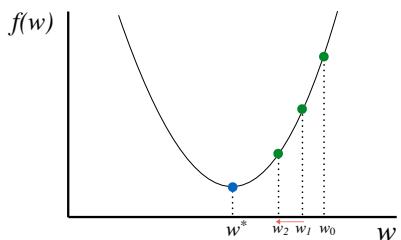
## Gradient Descent

Start at a random point

**Repeat**

- Determine a descent direction
- Choose a step size
- Update**

**Until** stopping criterion is satisfied



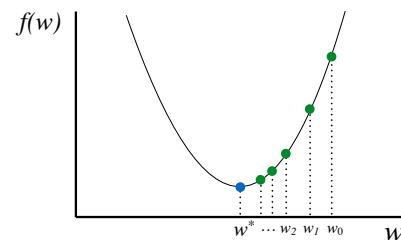
## Gradient Descent

Start at a random point

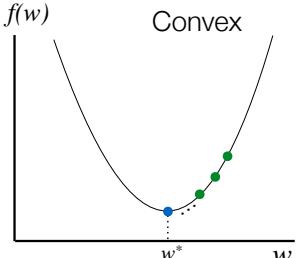
**Repeat**

- Determine a descent direction
- Choose a step size
- Update

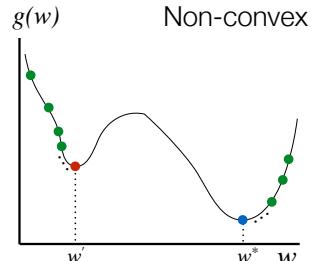
**Until** stopping criterion is satisfied



## Where Will We Converge?



Any local minimum is a global minimum



Multiple local minima may exist

## Gradient Descent:

**Batch gradient:** use error  $E_D(\mathbf{w})$  over entire training set  $D$

Do until satisfied:

1. Compute the gradient  $\nabla E_D(\mathbf{w}) = \left[ \frac{\partial E_D(\mathbf{w})}{\partial w_0} \dots \frac{\partial E_D(\mathbf{w})}{\partial w_n} \right]$
2. Update the vector of parameters:  $\mathbf{w} \leftarrow \mathbf{w} - \eta \nabla E_D(\mathbf{w})$

**Stochastic gradient:** use error  $E_d(\mathbf{w})$  over single examples  $d \in D$

Do until satisfied:

1. Choose (with replacement) a random training example  $d \in D$
2. Compute the gradient just for  $d$ :  $\nabla E_d(\mathbf{w}) = \left[ \frac{\partial E_d(\mathbf{w})}{\partial w_0} \dots \frac{\partial E_d(\mathbf{w})}{\partial w_n} \right]$
3. Update the vector of parameters:  $\mathbf{w} \leftarrow \mathbf{w} - \eta \nabla E_d(\mathbf{w})$

Stochastic approximates Batch arbitrarily closely as  
Stochastic can be much faster when  $D$  is very large  $\eta \rightarrow 0$   
Intermediate approach: use error over subsets of  $D$

## Maximize Conditional Log Likelihood: Gradient Ascent

$$\begin{aligned} l(W) &\equiv \ln \prod_l P(Y^l | X^l, W) = \sum_l \ln P(Y^l | X^l, W) \\ &= \sum_l Y^l (w_0 + \sum_i^n w_i X_i^l) - \ln(1 + \exp(w_0 + \sum_i^n w_i X_i^l)) \end{aligned}$$

$$\frac{\partial l(W)}{\partial w_i} = \sum_l X_i^l (Y^l - \hat{P}(Y^l = 1 | X^l, W))$$

Intuitive notion

## Maximize Conditional Log Likelihood: Gradient Ascent

$$\begin{aligned} l(W) &\equiv \ln \prod_l P(Y^l | X^l, W) = \sum_l \ln P(Y^l | X^l, W) \\ &= \sum_l Y^l (w_0 + \sum_i^n w_i X_i^l) - \ln(1 + \exp(w_0 + \sum_i^n w_i X_i^l)) \end{aligned}$$

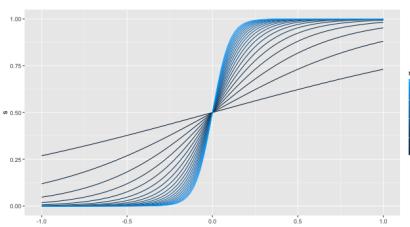
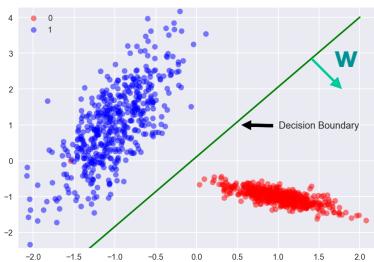
$$\frac{\partial l(W)}{\partial w_i} = \sum_l X_i^l (Y^l - \hat{P}(Y^l = 1 | X^l, W))$$

Gradient ascent algorithm: iterate until change  $< \varepsilon$   
For all  $i$ , repeat

$$w_i \leftarrow w_i + \eta \sum_l X_i^l (Y^l - \hat{P}(Y^l = 1 | X^l, W))$$

## Need to regularize the weights

- $w \rightarrow \infty$  to maximize the probability of the data, if data linearly separable



## That's all for M(C)LE. How about MAP?

- For MAP, need to define prior on W
  - given  $W = (w_1, \dots, w_n)$
  - let's assume prior  $P(w_i) = N(0, \sigma)$
  - i.e., assume zero mean, Gaussian prior for each  $w_i$
- A kind of Occam's razor (simplest is best) prior
- Helps avoid very large weights and overfitting

## MAP Estimates for Logistic Regression

$$W^{MAP} = \arg \max_W P(W|Y, X) = \frac{P(Y|W, X)P(W, X)}{P(Y, X)}$$

## MAP Estimates for Logistic Regression

$$W^{MAP} = \arg \max_W P(W|Y, X) = \frac{P(Y|W, X)P(W, X)}{P(Y, X)}$$

$$= \arg \max_W P(Y|W, X)P(W, X)$$

$$= \arg \max_W P(Y|W, X)P(W)P(X)$$

$$= \arg \max_W P(Y|W, X)P(W)$$

$$W^{MAP} = \arg \max_W [\ln P(Y|W, X) + \ln P(W)]$$

let's assume  
 $P(W, X) = P(W)P(X)$

zero mean  
Gaussian  
 $P(W)$

$$P(W) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{1}{2\sigma^2} \sum_i w_i^2\right)$$

$$W^{MAP} = \arg \max_W [\ln P(Y|W, X) - \left(\frac{1}{2\sigma^2} \sum_i w_i^2\right)]$$

## MLE vs MAP

- Maximum conditional likelihood estimate

$$W \leftarrow \arg \max_W \ln \prod_l P(Y^l | X^l, W)$$

$$w_i \leftarrow w_i + \eta \sum_l X_i^l (Y^l - \hat{P}(Y^l = 1 | X^l, W))$$

- Maximum a posteriori estimate with prior

$$W \leftarrow \arg \max_W \ln [P(W) \prod_l P(Y^l | X^l, W)]$$

$$w_i \leftarrow w_i - \eta \lambda w_i + \eta \sum_l X_i^l (Y^l - \hat{P}(Y^l = 1 | X^l, W))$$

## MAP estimates and Regularization

- Maximum a posteriori estimate with prior

$$W \leftarrow \arg \max_W \ln [P(W) \prod_l P(Y^l | X^l, W)]$$

$$w_i \leftarrow w_i - \eta \lambda w_i + \eta \sum_l X_i^l (Y^l - \hat{P}(Y^l = 1 | X^l, W))$$

called a “regularization” term

- helps reduce overfitting
- if  $P(W)$  is Gaussian, then encourages  $W$  to be near the mean of  $P(W)$ : zero here, but can easily use any mean
- used very frequently in Logistic Regression

## The Bottom Line

- Consider learning  $f: X \rightarrow Y$ , where
  - $X$  is a vector of real-valued features, ( $X_1 \dots X_n$ )
  - $Y$  is boolean
  - assume all  $X_i$  are conditionally independent given  $Y$
  - model  $P(X_i | Y = y_k)$  as Gaussian  $N(\mu_{ik}, \sigma_i)$
  - model  $P(Y)$  as Bernoulli ( $\pi$ )
- Then  $P(Y|X)$  is of this form, and we can directly estimate  $W$ 
$$P(Y = 1 | X = (X_1, \dots, X_n)) = \frac{1}{1 + \exp(w_0 + \sum_i w_i X_i)}$$
- Furthermore, same holds if the  $X_i$  are boolean
  - trying proving that to yourself

## Generative vs. Discriminative Classifiers

Training classifiers involves estimating  $f: X \rightarrow Y$ , or  $P(Y|X)$

Generative classifiers (e.g., Naïve Bayes)

- Assume some functional form for  $P(X|Y)$ ,  $P(X)$
- Estimate parameters of  $P(X|Y)$ ,  $P(X)$  directly from training data
- Use Bayes rule to calculate  $P(Y|X=x)$

Discriminative classifiers (e.g., Logistic regression)

- Assume some functional form for  $P(Y|X)$
- Estimate parameters of  $P(Y|X)$  directly from training data

## Use Naïve Bayes or Logistic Regression?

Consider

- Restrictiveness of modeling assumptions
- Rate of convergence (in amount of training data) toward asymptotic hypothesis

## Naïve Bayes vs Logistic Regression

Consider Y boolean,  $X_i$  continuous,  $X=(X_1 \dots X_n)$

Number of parameters to estimate:

- NB:

$$P(x | y_k) = \frac{1}{\sigma_{ik}\sqrt{2\pi}} e^{-\frac{(x-\mu_{ik})^2}{2\sigma_{ik}^2}}$$

- LR:

$$P(Y = 0|X, W) = \frac{1}{1 + \exp(w_0 + \sum_i w_i X_i)}$$

$$P(Y = 1|X, W) = \frac{\exp(w_0 + \sum_i w_i X_i)}{1 + \exp(w_0 + \sum_i w_i X_i)}$$

## Naïve Bayes vs Logistic Regression

Consider Y boolean,  $X_i$  continuous,  $X=(X_1 \dots X_n)$

Number of parameters:

- NB:  $4n + 1$
- LR:  $n+1$

Estimation method:

- NB parameter estimates are uncoupled
- LR parameter estimates are coupled

## G.Naïve Bayes vs. Logistic Regression

Recall two assumptions deriving form of LR from GNBayes:

1.  $X_i$  conditionally independent of  $X_k$  given Y
2.  $P(X_i | Y = y_k) = N(\mu_{ik}, \sigma_i)$ ,  $\leftarrow$  not  $N(\mu_{ik}, \sigma_{ik})$

[Ng & Jordan, 2002]

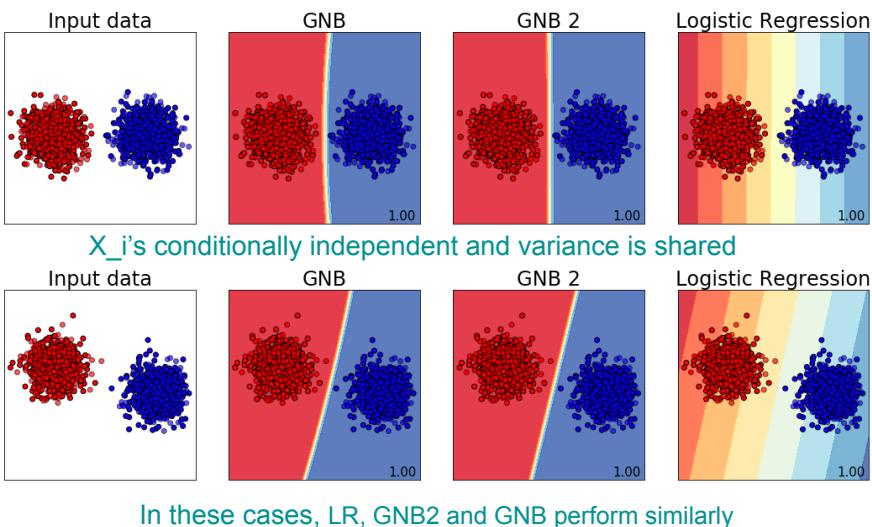
Consider three learning methods:

- GNB (assumption 1 only) -- decision surface can be non-linear
- GNB2 (assumption 1 and 2) -- decision surface linear
- LR -- decision surface linear, trained without assumption 1.

How do these methods perform if we have plenty of data and:

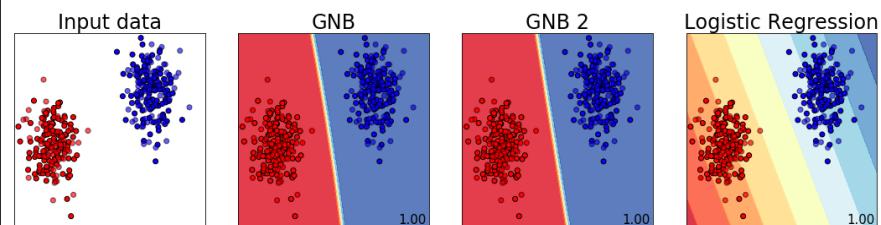
- Both (1) and (2) are satisfied:

**Assumptions 1 and 2 are satisfied**



## Assumptions 1 and 2 satisfied

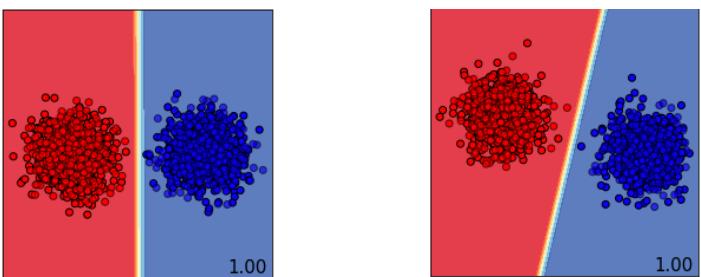
The decision boundary of GNB and GNB2 is sensitive to the locations of the means (since the variances are the same)



## Assumptions 1 and 2 satisfied

If the variances of the  $X_i$  are the same (across classes and across  $i$ ), the decision boundary of GNB2 and GNB is determined by the distance to the mean (perpendicular bisector)

If one of the coordinates of the two means are the same, then the decision boundary becomes perpendicular to that axis



## G.Naïve Bayes vs. Logistic Regression

Recall two assumptions deriving form of LR from GNBayes:

1.  $X_i$  conditionally independent of  $X_k$  given  $Y$
  2.  $P(X_i | Y = y_k) = N(\mu_{ik}, \sigma_i)$ ,  $\leftarrow$  not  $N(\mu_{ik}, \sigma_{ik})$

[Ng & Jordan, 2002]

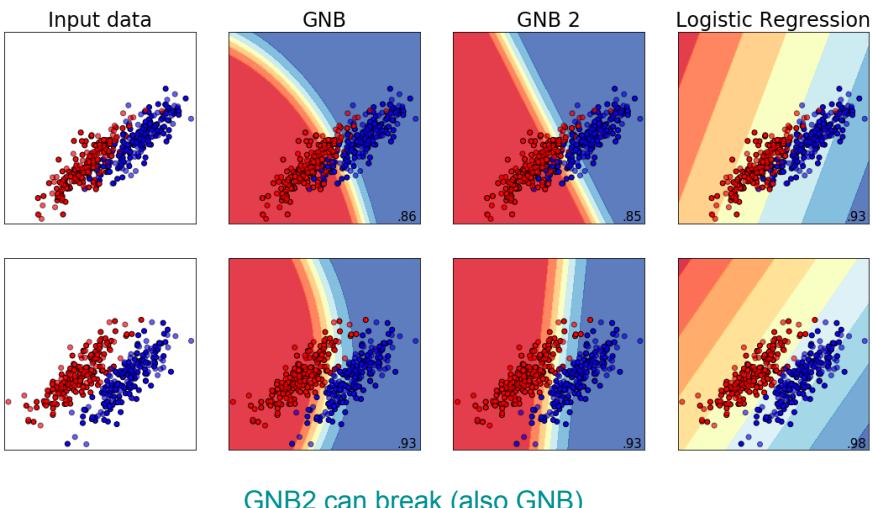
Consider three learning methods:

- GNB (assumption 1 only) -- decision surface can be non-linear
  - GNB2 (assumption 1 and 2) -- decision surface linear
  - LR -- decision surface linear, trained without assumption 1.

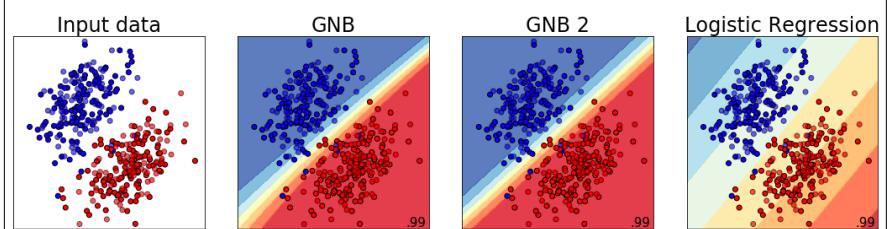
How do these methods perform if we have plenty of data and:

- Both (1) and (2) are satisfied
  - (2) is satisfied, but not (1)

## Assumption 2 satisfied and not 1

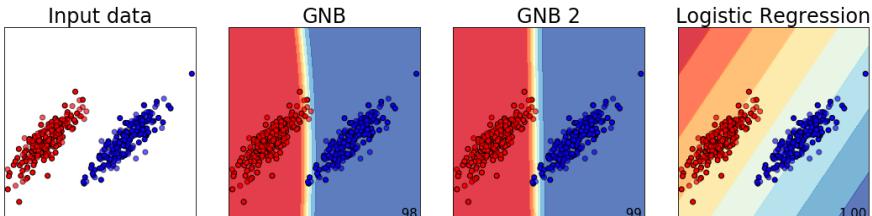


## Assumption 2 satisfied and not 1



GNB2 and GNB can also work well

## Assumption 2 satisfied and not 1



The decision boundary of GNB2 and GNB is also dependent on the means of the two classes. If one of the coordinates of the two means is the same, again, we have a decision boundary parallel to that axis

# G.Naïve Bayes vs. Logistic Regression

Recall two assumptions deriving form of LR from GNBayes:

1.  $X_i$  conditionally independent of  $X_k$  given  $Y$
  2.  $P(X_i | Y = y_k) = N(\mu_{ik}, \sigma_i)$ ,  $\leftarrow$  not  $N(\mu_{ik}, \sigma_{ik})$

[Ng & Jordan, 2002]

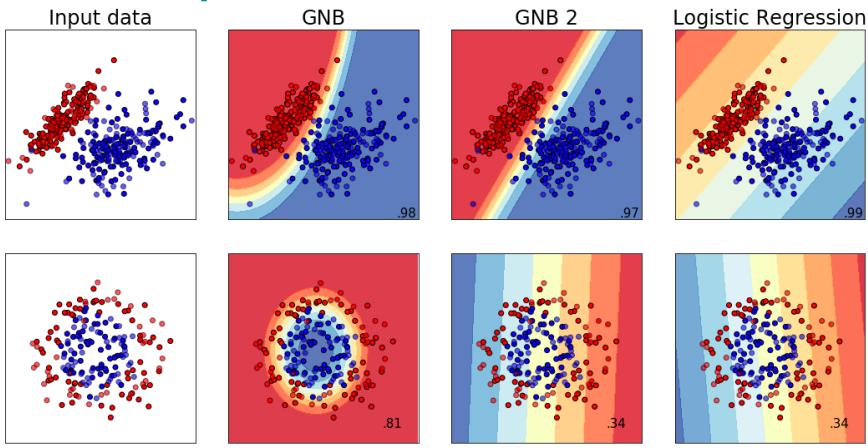
Consider three learning methods:

- GNB (assumption 1 only) -- decision surface can be non-linear
  - GNB2 (assumption 1 and 2) – decision surface linear
  - LR -- decision surface linear, trained without assumption 1.

How do these methods perform if we have plenty of data and:

- Both (1) and (2) are satisfied
  - (2) is satisfied, but not (1)
  - Neither (1) nor (2) is satisfied

## Assumptions 1 and 2 are not satisfied



Depending on the dataset, GNB and LR have different performances.  
Even though LR and GNB2 can be expressed in the same way, LR has more flexibility to learn parameters that fit the data, and they are don't have to be tied to the marginal means and variance

## G.Naïve Bayes vs. Logistic Regression

What if we have only finite training data?

They converge at different rates to their asymptotic ( $\infty$  data) error

Let  $\epsilon_{A,m}$  refer to expected error of learning algorithm A after  $m$  training examples

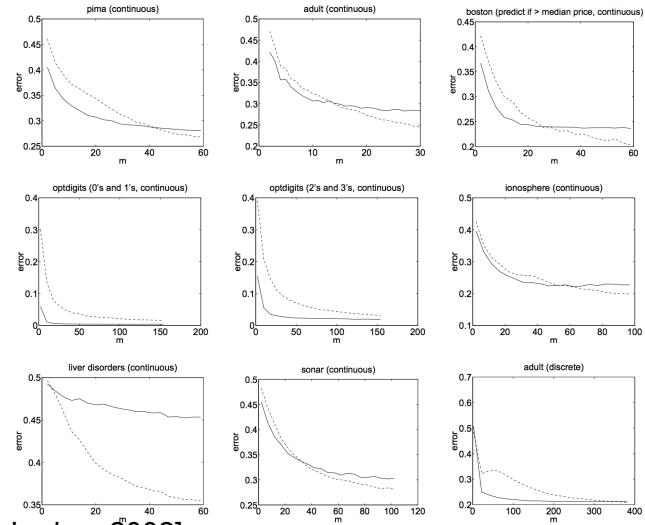
Let  $n$  be the number of features:  $(X_1 \dots X_n)$  [Ng & Jordan, 2002]

$$\epsilon_{LR,m} \leq \epsilon_{LR,\infty} + O\left(\sqrt{\frac{n}{m}}\right)$$

$$\epsilon_{GNB,m} \leq \epsilon_{GNB,\infty} + O\left(\sqrt{\frac{\log(n)}{m}}\right)$$

So, GNB requires  $m = O(\log n)$  to converge, but LR requires  $m = O(n)$

## Some experiments from UCI data sets



[Ng & Jordan, 2002]

## Some experiments from UCI data sets

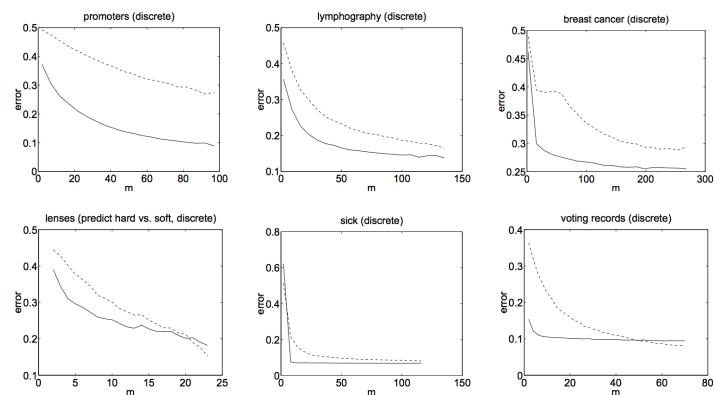


Figure 1: Results of 15 experiments on datasets from the UCI Machine Learning repository. Plots are of generalization error vs.  $m$  (averaged over 1000 random train/test splits). Dashed line is logistic regression; solid line is naïve Bayes.

[Ng & Jordan, 2002]

## Naïve Bayes vs. Logistic Regression

The bottom line:

GNB2 and LR both use linear decision surfaces, GNB need not

Given infinite data, LR is better or equal to GNB2 because *training procedure* does not make assumptions 1 or 2 (though our derivation of the form of  $P(Y|X)$  did).

But GNB2 converges more quickly to its perhaps-less-accurate asymptotic error. (more bias than LR)

And GNB is both more biased (assumption1) and less (no assumption 2) than LR, so either might outperform the other.

## What you should know:

- Logistic regression
  - Functional form follows from Naïve Bayes assumptions
    - For Gaussian Naïve Bayes assuming variance  $\sigma_{i,k} = \sigma_i$
    - For discrete-valued Naïve Bayes too
  - But training procedure picks parameters without making conditional independence assumption
  - MLE training: pick  $W$  to maximize  $P(Y | X, W)$
  - MAP training: pick  $W$  to maximize  $P(W | X, Y)$ 
    - ‘regularization’
    - helps reduce overfitting
- Gradient ascent/descent
  - General approach when closed-form solutions unavailable
- Generative vs. Discriminative classifiers
  - Bias vs. variance tradeoff

## Questions to think about:

- Can you use Naïve Bayes for a combination of discrete and real-valued  $X_i$ ?
- How can we easily model the assumption that just 2 of the  $n$  attributes are dependent?
- What does the decision surface of a Naïve Bayes classifier look like?
- How would you select a subset of  $X_i$ 's?